

# What Would be included...

## • 📞 Telecom Customer Churn Prediction

This project predicts whether a telecom customer will **churn (leave the company)** or **stay**, using machine learning models such as Logistic Regression and Random Forest. The goal is to help telecom companies **reduce customer loss** by identifying customers likely to churn.

---

### Dataset Overview

The dataset contains various information about telecom customers, including demographics, account details, and subscription types.

Column Name	Description
customerID	Unique ID assigned to each customer
Gender	Gender of the customer (Male/Female)
SeniorCitizen	Indicates if the customer is a senior citizen (1 = Yes, 0 = No)
Partner	Whether the customer has a partner (Yes/No)
Dependents	Whether the customer has dependents (Yes/No)
Tenure	Number of months the customer has stayed with the company
PhoneService	Whether the customer has a phone service (Yes/No)
MultipleLines	Whether the customer has multiple lines (Yes/No/No phone service)
InternetService	Type of internet service (DSL, Fiber optic, No)
OnlineSecurity	Whether the customer has online security (Yes/No/No internet service)
DeviceProtection	Whether the customer has device protection (Yes/No/No internet service)
TechSupport	Whether the customer has tech support (Yes/No/No internet service)
StreamingTV	Whether the customer has streaming TV service (Yes/No/No internet service)
StreamingMovies	Whether the customer has streaming movies (Yes/No/No internet service)
Contract	Type of contract (Month-to-month, One year, Two year)
PaperlessBilling	Whether billing is paperless (Yes/No)
PaymentMethod	Method of payment (e.g., Electronic check, Credit card, etc.)

Column Name	Description
MonthlyCharges	Amount charged per month
TotalCharges	Total amount charged to the customer
Churn	Target variable: whether the customer churned (Yes/No)


## Objective

Predict whether a customer will churn based on their account and service details.

## Author

Name	Ali Husnain
Roll No	SU92-BSAIM-F24-024

In [160...

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.ensemble import RandomForestClassifier #  <-- Add this import
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```


In [161...

```
df = pd.read_csv('TelecomCustomerChurn.csv')
df
```

Out[161...

	customerID	Gender	SeniorCitizen	Partner	Dependents	Tenure	PhoneService	Mu
0	7590-VHVEG	Female	0	Yes	No	1	No	
1	5575-GNVDE	Male	0	No	No	34	Yes	
2	3668-QPYBK	Male	0	No	No	2	Yes	
3	7795-CFOCW	Male	0	No	No	45	No	
4	9237-HQITU	Female	0	No	No	2	Yes	
...	...	...	...	...	...	...	...	...
7038	6840-RESVB	Male	0	Yes	Yes	24	Yes	
7039	2234-XADUH	Female	0	Yes	Yes	72	Yes	
7040	4801-JAZZL	Female	0	Yes	Yes	11	No	
7041	8361-LTMKD	Male	1	Yes	No	4	Yes	
7042	3186-AJIEK	Male	0	No	No	66	Yes	

7043 rows × 21 columns



In [162...

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   object
1   Gender                 7043 non-null   object
2   SeniorCitizen          7043 non-null   int64
3   Partner                7043 non-null   object
4   Dependents             7043 non-null   object
5   Tenure                 7043 non-null   int64
6   PhoneService           7043 non-null   object
7   MultipleLines           7043 non-null   object
8   InternetService        7043 non-null   object
9   OnlineSecurity         7043 non-null   object
10  OnlineBackup           7043 non-null   object
11  DeviceProtection       7043 non-null   object
12  TechSupport            7043 non-null   object
13  StreamingTV            7043 non-null   object
14  StreamingMovies        7043 non-null   object
15  Contract               7043 non-null   object
16  PaperlessBilling       7043 non-null   object
17  PaymentMethod          7043 non-null   object
18  MonthlyCharges         7043 non-null   float64
19  TotalCharges           7043 non-null   object
20  Churn                  7043 non-null   object
dtypes: float64(1), int64(2), object(18)
memory usage: 1.1+ MB
```

```
In [163... df.describe()
```

Out[163...

	SeniorCitizen	Tenure	MonthlyCharges
count	7043.000000	7043.000000	7043.000000
mean	0.162147	32.371149	64.761692
std	0.368612	24.559481	30.090047
min	0.000000	0.000000	18.250000
25%	0.000000	9.000000	35.500000
50%	0.000000	29.000000	70.350000
75%	0.000000	55.000000	89.850000
max	1.000000	72.000000	118.750000

```
In [164... df.isnull()
```

Out[164...

	customerID	Gender	SeniorCitizen	Partner	Dependents	Tenure	PhoneService	Mu
0	False	False	False	False	False	False	False	
1	False	False	False	False	False	False	False	
2	False	False	False	False	False	False	False	
3	False	False	False	False	False	False	False	
4	False	False	False	False	False	False	False	
...	...	...	...	...	...	...	...	...
7038	False	False	False	False	False	False	False	
7039	False	False	False	False	False	False	False	
7040	False	False	False	False	False	False	False	
7041	False	False	False	False	False	False	False	
7042	False	False	False	False	False	False	False	

7043 rows × 21 columns



In [165...

df.isnull().sum()

Out[165...

```
customerID      0
Gender          0
SeniorCitizen   0
Partner         0
Dependents      0
Tenure          0
PhoneService    0
MultipleLines   0
InternetService 0
OnlineSecurity  0
OnlineBackup    0
DeviceProtection 0
TechSupport     0
StreamingTV     0
StreamingMovies 0
Contract        0
PaperlessBilling 0
PaymentMethod   0
MonthlyCharges  0
TotalCharges    0
Churn           0
dtype: int64
```

- We are Clearly See there is no null values

In [166...

df.duplicated().sum()


Out[166...] np.int64(0)

In [167...] `df.columns`

Out[167...] Index(['customerID', 'Gender', 'SeniorCitizen', 'Partner', 'Dependents', 'Tenure', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'], dtype='object')


In [168...] `sns.heatmap(df.isnull(), cbar=False)`

Out[168...] <Axes: >



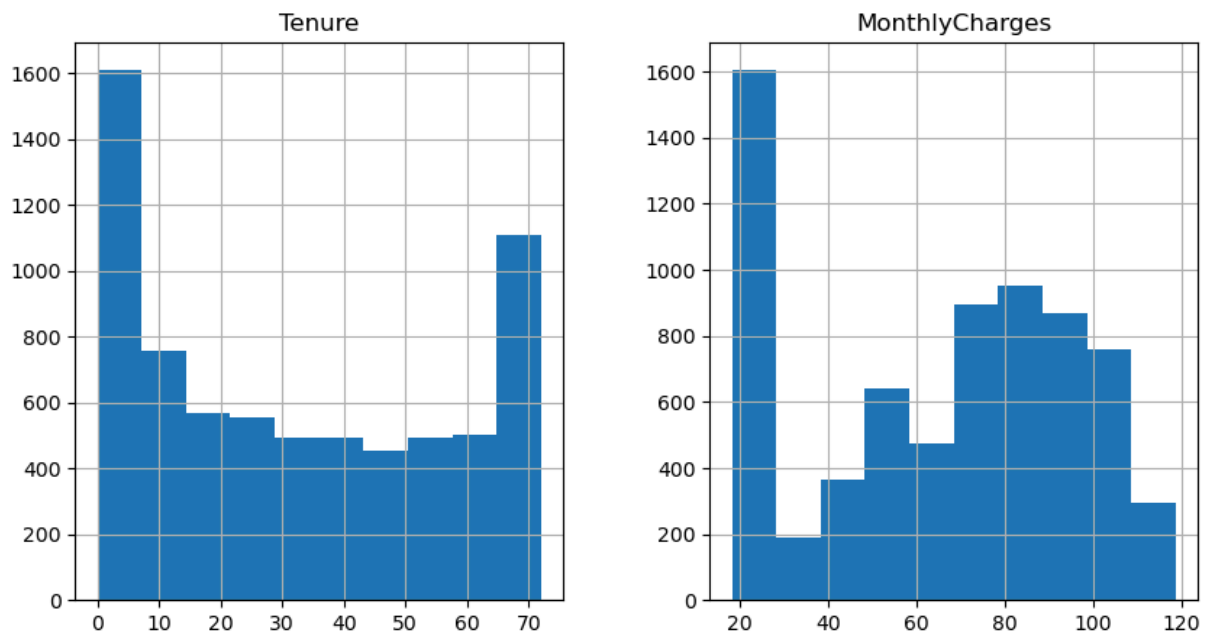
In [169...] `df['Churn'].value_counts(normalize=True)`  
`sns.countplot(x='Churn', data=df)`

Out[169...] <Axes: xlabel='Churn', ylabel='count'>



```
In [170...] num_cols = ['Tenure', 'MonthlyCharges', 'TotalCharges']
df[num_cols].hist(figsize=(10, 5))
```

```
Out[170...] array([[<Axes: title={'center': 'Tenure'}>,
      <Axes: title={'center': 'MonthlyCharges'}>]], dtype=object)
```



```
In [171...] df.drop("customerID", axis=1, inplace=True)
```

```
In [172...] df["TotalCharges"] = pd.to_numeric(df["TotalCharges"], errors="coerce")
```

```
df["TotalCharges"].fillna(df["TotalCharges"].mean(), inplace=True)
```

C:\Users\alihu\AppData\Local\Temp\ipykernel\_17376\3639554342.py:2: FutureWarning: A value is trying to be set on a copy of a DataFrame or Series through chained assignment using an inplace method.

The behavior will change in pandas 3.0. This inplace method will never work because the intermediate object on which we are setting values always behaves as a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.


```
df["TotalCharges"].fillna(df["TotalCharges"].mean(), inplace=True)
```

```
In [173... X = df.drop("Churn", axis=1)
y = df["Churn"]
```

```
In [174... le = LabelEncoder()
for col in X.select_dtypes(include=['object']).columns:
    X[col] = le.fit_transform(X[col])
y = le.fit_transform(y)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

```
In [175... model = RandomForestClassifier(random_state=42)
model.fit(X_train, y_train)
importances = pd.Series(model.feature_importances_, index=X.columns)
importances.nlargest(10).plot(kind='barh')
```

Out[175... <Axes: >





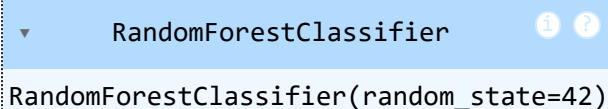
In [176... `X_train.dtypes`

```
Out[176... Gender                int64
SeniorCitizen          int64
Partner                int64
Dependents              int64
Tenure                  int64
PhoneService            int64
MultipleLines           int64
InternetService         int64
OnlineSecurity           int64
OnlineBackup            int64
DeviceProtection        int64
TechSupport             int64
StreamingTV             int64
StreamingMovies          int64
Contract                int64
PaperlessBilling         int64
PaymentMethod           int64
MonthlyCharges          float64
TotalCharges            float64
dtype: object
```

In [177... `feature_names = X.columns.tolist()`

In [178... `scaler = StandardScaler()`  
`X_train = scaler.fit_transform(X_train)`  
`X_test = scaler.transform(X_test)`

In [179... `model = RandomForestClassifier(random_state=42)`  
`model.fit(X_train, y_train)`

Out[179...  `RandomForestClassifier`  
`RandomForestClassifier(random_state=42)`

In [180... `y_pred = model.predict(X_test)`  
`print("✅ Accuracy:", accuracy_score(y_test, y_pred))`  
`print("📊 Confusion Matrix:\n", confusion_matrix(y_test, y_pred))`  
`print("📄 Classification Report:\n", classification_report(y_test, y_pred))`

✓ Accuracy: 0.791292001893043

📊 Confusion Matrix:

[[1401 138]

[ 303 271]]

📄 Classification Report:

	precision	recall	f1-score	support
0	0.82	0.91	0.86	1539
1	0.66	0.47	0.55	574
accuracy			0.79	2113
macro avg	0.74	0.69	0.71	2113
weighted avg	0.78	0.79	0.78	2113

In [182...

```
import pickle
model_data = {
    "model": model,
    "scaler": scaler,
    "feature_names": feature_names
}

with open("model.pkl", "wb") as file:
    pickle.dump(model_data, file)

print("✓ Model file saved successfully with feature names!")
```

✓ Model file saved successfully with feature names!

In [ ]: